

# 【OS】 Day7(3)

## 【Ch2】 Process Homework(Simulation)

### Exercise 1

1. Run `process-run.py` with the following flags: `-l 5:100,5:100`. What should the CPU utilization be (e.g., the percent of time the CPU is in use?) Why do you know this? Use the `-c` and `-p` flags to see if you were right.

The CPU executes `process0` first. When it finishes executing `process0`, it will start to execute `process1`. The total time cost is 10 time units.

```
PS D:\ostep-homework\cpu-intro> python .\process-run.py -l 5:100,5:100 -c -p
Time    PID: 0    PID: 1    CPU    I/Os
1       RUN:cpu  READY    1
2       RUN:cpu  READY    1
3       RUN:cpu  READY    1
4       RUN:cpu  READY    1
5       RUN:cpu  READY    1
6       DONE    RUN:cpu  1
7       DONE    RUN:cpu  1
8       DONE    RUN:cpu  1
9       DONE    RUN:cpu  1
10      DONE    RUN:cpu  1

Stats: Total Time 10
Stats: CPU Busy 10 (100.00%)
Stats: IO Busy 0 (0.00%)
```

### Exercise 2

2. Now run with these flags: `./process-run.py -l 4:100,1:0`. These flags specify one process with 4 instructions (all to use the CPU), and one that simply issues an I/O and waits for it to be done. How long does it take to complete both processes? Use `-c` and `-p` to find out if you were right.

CPU will **execute process0 first**, which has 4 CPU instructions and cost 4 time units.

**After finishing process0**, it will execute process1, which have one I/O instruction(cost 2 time units) and keep the I/O busy for 5 time units.

Thus, the execution of the two processes will cost 11 time units.

```
PS D:\ostep-homework\cpu-intro> python .\process-run.py -l 4:100,1:0 -c -p
Time   PID: 0      PID: 1      CPU      IOs
1      RUN:cpu     READY      1
2      RUN:cpu     READY      1
3      RUN:cpu     READY      1
4      RUN:cpu     READY      1
5      DONE      RUN:io      1
6      DONE      WAITING     1
7      DONE      WAITING     1
8      DONE      WAITING     1
9      DONE      WAITING     1
10     DONE      WAITING     1
11*    DONE      RUN:io_done 1
Stats: Total Time 11
Stats: CPU Busy 6 (54.55%)
Stats: IO Busy 5 (45.45%)
```

### Exercise 3

3. Switch the order of the processes: `-l 1:0, 4:100`. What happens now? Does switching the order matter? Why? (As always, use `-c` and `-p` to see if you were right)

*Switch order matters.*

In this case, the I/O instruction is executed first. CPU **can run process1 when process0 is waiting for I/O to be done**.

After process1 is finished, there is still 1 time unit left for the I/O to do its job and 1 time unit cost for the CPU to end the I/O.

Thus, the total time cost is 7 time units.

### Exercise 4

- We'll now explore some of the other flags. One important flag is `-S`, which determines how the system reacts when a process issues an I/O. With the flag set to `SWITCH_ON_END`, the system will NOT switch to another process while one is doing I/O, instead waiting until the process is completely finished. What happens when you run the following two processes (`-1 1:0,4:100 -c -S SWITCH_ON_END`), one doing I/O and the other doing CPU work?

With the `-S` flag set to `SWITCH_ON_END`, the CPU is **not doing work when the I/O instruction is being executed**.

```
(otherwise states are not printed)
PS C:\ostep-homework\cpu-intro> python .\process-run.py -1 1:0,4:100 -c -S SWITCH_ON_END
Time      PID: 0      PID: 1      CPU      I/Os
1         RUN:io    READY      1
2         WAITING  READY      1
3         WAITING  READY      1
4         WAITING  READY      1
5         WAITING  READY      1
6         WAITING  READY      1
7*        RUN:io_done  READY      1
8         DONE     RUN:cpu     1
9         DONE     RUN:cpu     1
10        DONE     RUN:cpu     1
11        DONE     RUN:cpu     1
```

### Exercise 5

- Now, run the same processes, but with the switching behavior set to switch to another process whenever one is WAITING for I/O (`-1 1:0,4:100 -c -S SWITCH_ON_IO`). What happens now? Use `-c` and `-p` to confirm that you are right.

Now, process1(the 4 CPU instructions) will be executed while waiting for I/O. T

he I/O instruction executes for 1 time unit, then **passes the control to process1 and have it execute for 4 time units**.

Then, another time unit is spent for I/O. Finally, the control was passed to CPU when the I/O is done.

### Exercise 6

6. One other important behavior is what to do when an I/O completes. With `-I IO_RUN_LATER`, when an I/O completes, the process that issued it is not necessarily run right away; rather, whatever was running at the time keeps running. What happens when you run this combination of processes? (Run `./process-run.py -l 3:0,5:100,5:100,5:100 -S SWITCH_ON_IO -I IO_RUN_LATER -c -p`) Are system resources being effectively utilized?

The system resources are not being effectively utilized. After the first I/O finishes, instead of executing the second I/O instruction and then the second CPU instruction, the second CPU instruction was executed.

This caused a waste of time later. During the second I/O instruction is being executed, the CPU just waits without doing any work.

```
PS C:\ostep-homework\cpu-intro> python .\process-run.py -l 3:0,5:100,5:100,5:100 -S SWITCH_ON_IO -I IO_RUN_LATER -c -p
Time   PID: 0      PID: 1      PID: 2      PID: 3      CPU      I/Os
1      RUN:io      READY      READY      READY      1
2      WAITING    RUN:cpu    READY      READY      1      1
3      WAITING    RUN:cpu    READY      READY      1      1
4      WAITING    RUN:cpu    READY      READY      1      1
5      WAITING    RUN:cpu    READY      READY      1      1
6      WAITING    RUN:cpu    READY      READY      1      1
7*     READY      DONE      RUN:cpu    READY      1
8      READY      DONE      RUN:cpu    READY      1
9      READY      DONE      RUN:cpu    READY      1
10     READY      DONE      RUN:cpu    READY      1
11     READY      DONE      RUN:cpu    READY      1
12     READY      DONE      DONE      RUN:cpu    1
13     READY      DONE      DONE      RUN:cpu    1
14     READY      DONE      DONE      RUN:cpu    1
15     READY      DONE      DONE      RUN:cpu    1
16     READY      DONE      DONE      RUN:cpu    1
17     RUN:io_done  DONE      DONE      DONE      1
18     RUN:io      DONE      DONE      DONE      1
19     WAITING    DONE      DONE      DONE      1      1
20     WAITING    DONE      DONE      DONE      1      1
21     WAITING    DONE      DONE      DONE      1      1
22     WAITING    DONE      DONE      DONE      1      1
23     WAITING    DONE      DONE      DONE      1      1
24*    RUN:io_done  DONE      DONE      DONE      1
25     RUN:io      DONE      DONE      DONE      1
26     WAITING    DONE      DONE      DONE      1      1
27     WAITING    DONE      DONE      DONE      1      1
28     WAITING    DONE      DONE      DONE      1      1
29     WAITING    DONE      DONE      DONE      1      1
30     WAITING    DONE      DONE      DONE      1      1
31*    RUN:io_done  DONE      DONE      DONE      1
Stats: Total Time 31
Stats: CPU Busy 21 (67.74%)
Stats: IO Busy 15 (48.39%)
```

## Exercise 7

- Now run the same processes, but with `-I IO_RUN_IMMEDIATE` set, which immediately runs the process that issued the I/O. How does this behavior differ? Why might running a process that just completed an I/O again be a good idea?

Now, the CPU is fully utilized.

```
PS C:\ostep-homework\cpu-intro> python .\process-run.py -l 3:0,5:100,5:100,5:100 -S SWITCH_ON_IO -I IO_RUN_IMMEDIATE -c
-p
Time      PID: 0      PID: 1      PID: 2      PID: 3      CPU      IOs
1         RUN:io    READY      READY      READY      1         1
2         WAITING  RUN:cpu    READY      READY      1         1
3         WAITING  RUN:cpu    READY      READY      1         1
4         WAITING  RUN:cpu    READY      READY      1         1
5         WAITING  RUN:cpu    READY      READY      1         1
6         WAITING  RUN:cpu    READY      READY      1         1
7*        RUN:io_done  DONE      READY      READY      1         1
8         RUN:io    DONE      READY      READY      1         1
9         WAITING  DONE      RUN:cpu    READY      1         1
10        WAITING  DONE      RUN:cpu    READY      1         1
11        WAITING  DONE      RUN:cpu    READY      1         1
12        WAITING  DONE      RUN:cpu    READY      1         1
13        WAITING  DONE      RUN:cpu    READY      1         1
14*       RUN:io_done  DONE      DONE      READY      1         1
15        RUN:io    DONE      DONE      READY      1         1
16        WAITING  DONE      DONE      RUN:cpu    1         1
17        WAITING  DONE      DONE      RUN:cpu    1         1
18        WAITING  DONE      DONE      RUN:cpu    1         1
19        WAITING  DONE      DONE      RUN:cpu    1         1
20        WAITING  DONE      DONE      RUN:cpu    1         1
21*       RUN:io_done  DONE      DONE      DONE      1         1

Stats: Total Time 21
Stats: CPU Busy 21 (100.00%)
Stats: IO Busy 15 (71.43%)
```

### Exercise 8

- Now run with some randomly generated processes: `-s 1 -l 3:50,3:50` or `-s 2 -l 3:50,3:50` or `-s 3 -l 3:50,3:50`. See if you can predict how the trace will turn out. What happens when you use the flag `-I IO_RUN_IMMEDIATE` vs. `-I IO_RUN_LATER`? What happens when you use `-S SWITCH_ON_IO` vs. `-S SWITCH_ON_END`?

When `IO_RUN_IMMEDIATE` is set, the CPU will run another I/O instruction right after the current instruction finishes, and thus fully utilizes the time waiting for the data to be fetched.

When `SWITCH_ON_IO` is set, the CPU transfers control to another process when an I/O instruction is invoked.