

# 【OS】 Day8(2)

## 【Ch2】 Process API(2)

### 5.3 The exec() System Call

The `exec()` system call is useful when you want to run a program that is different from the calling program.

For example, calling `fork()` in `p2.c` is only useful if you want to keep running copies of the same program.

However, often you want to run a different program; `exec()` does just that.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("Hello world (pid: %d)\n", (int)getpid());
    int rc = fork();
    if(rc < 0) {
        fprintf(stderr, "Fork failed\n");
        exit(1);
    } else if(rc == 0) {
        printf("Hello, I am child (pid: %d)\n", (int)getpid());
        char *args[3];
        args[0] = strdup("wc");
        args[1] = strdup("prac.c");
        args[2] = NULL;
        execvp(args[0], args);
    } else {
        int rc_wait = wait(NULL);
        printf("hello, I am parent of %d (wait: %d) (pid: %d)\n", rc, rc_wait, (int)getpid());
    }
    return 0;
}
```

In this example, the child process calls `execvp()` in order to run the program `wc`, which is the word counting program. In fact, it runs `wc` on the source file `p3.c`, thus telling us how many lines, words, and bytes are found in the file:

```

prompt> ./p3
hello world (pid:29383)
hello, I am child (pid:29384)
      29      107      1030 p3.c
hello, I am parent of 29384 (rc_wait:29384) (pid:29383)
prompt>

```

What `exec()` does: given the name of an executable(e.g. `wc`), and some arguments(e.g. `p3.c`) it loads code (and static data) from that executable and overwrites it current code segment (and current static data) with it; the heap and stack and other parts of the memory space of the program are re-initialized.

Thus, the OS simply runs that program, passing in any arguments as the `argv` of that process.

Thus, it does not create a new process; rather, it transforms the currently running program(formerly `p3`) into a different running program(`wc`). After the `exec()` in the child, it is almost as if `p3.c` never ran; a successful call to `exec()` never returns.

## 5.4 Why? Motivating The API

*Why would we build such an odd interface to what should be the simple act of creating a new process?*

The separation of `fork()` and `exec()` is essential in building a UNIX shell, because it lets the shell run code after the call to `fork()` but before the call to `exec()`; this code can alter the environment of the about-to-be-run program, and thus enables a variety of interesting features to be readily built.

The shell is just a user program. It shows you a prompt and then wait for you to type something into it.

You then type a command(i.e. the name of an executable program, plus any arguments) into it; in most cases, the shell then figures out where in the file system the executable resides, calls `fork()` to create a new child process to run the command, calls some variant of `exec()` to run the command, and then waits for the command to complete by calling `wait()`.

When the child completes, the shell returns from `wait()` and prints out a prompt again, ready for your next command.

Let's look at an example:

```
prompt> wc p3.c > newfile.txt
```

In the example above, the output of the program `wc` is redirected into the output file `newfile.txt` (the greater-than sign is how said redirection is indicated).

The way the shell accomplishes this task is quite simple: when the child is created, before calling `exec()`, the shell closes standard output and opens the file `newfile.txt`. By doing so, any output from the soon-to-be-running program `wc` are sent to the file instead of the screen.

The following program does exactly this:

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <unistd.h>
4  #include <string.h>
5  #include <fcntl.h>
6  #include <sys/wait.h>
7
8  int main(int argc, char *argv[]) {
9      int rc = fork();
10     if (rc < 0) {
11         // fork failed
12         fprintf(stderr, "fork failed\n");
13         exit(1);
14     } else if (rc == 0) {
15         // child: redirect standard output to a file
16         close(STDOUT_FILENO);
17         open("./p4.output", O_CREAT|O_WRONLY|O_TRUNC, S_IRWXU);
18
19         // now exec "wc"...
20         char *myargs[3];
21         myargs[0] = strdup("wc"); // program: wc (word count)
22         myargs[1] = strdup("p4.c"); // arg: file to count
23         myargs[2] = NULL; // mark end of array
24         execvp(myargs[0], myargs); // runs word count
25     } else {
26         // parent goes down this path (main)
27         int rc_wait = wait(NULL);
28     }
29     return 0;
30 }

```

Figure 5.4: All Of The Above With Redirection (p4.c)

The reason this redirection works is due to an assumption about how the operating system manages file descriptors.

Specifically, UNIX systems **start looking for free file descriptors at zero**. In this case, `STDOUT_FILENO` will be the first available one and thus get assigned when `open()` is called.

Subsequent **writes by the child process to the standard output file descriptor**, for example by routines such as `printf()`, will then be **routed transparently to the newly-opened file instead of the screen**.

Here is the output:

```
prompt> ./p4
prompt> cat p4.output
      32      109      846 p4.c
prompt>
```