

【OS】 Day15(2)

【Ch9】 Scheduling: Proportional Share Scheduling

9.7 The Linux Completely Fair Scheduling(2)

Using Red-Black Trees

For a scheduler, there are many facets of efficiency, but one of them is as simple as this: when the scheduler has **to find the next job to run**, it should do so as quickly as possible.

CFS addresses this by keeping processes in a **red-black tree**.

CFS does not keep all process in this structure rather, **only running(or runnable) processes are kept therein**. If a process goes to sleep(say, waiting on an I/O to complete), it is removed from the tree and kept track of elsewhere.

Keeping the same values in a red-black tree makes most operations more efficient. Processes are **ordered in the tree by vruntime**, and most operations are logarithmic in time($O(\log n)$). When n is in the thousands, logarithmic is noticeably more efficient than linear.




Figure 9.5: CFS Red-Black Tree

Dealing With I/O And Sleeping Processes

One problem with picking the lowest `vruntime` is to run next arises with jobs that **have gone to sleep for a long period of time**. Imagine we have a process who has slept a long period of time, when it wakes up, its `vruntime` will be **far behind other processes**, and thus will monopolize the CPU.

CFS handles this case by altering the `vruntime` of a job when it wakes up. Specifically, CFS **sets the `vruntime` of that job to the minimum value** found in the tree. In this way, CFS avoids starvation, but not without a cost: jobs that sleep for short periods of time frequently do not ever get their fair share of the CPU.