

# 【OS】 Day32(2)

▼ Class	Operating System: Three Easy Pieces
📅 Date	@February 5, 2022

## 【Ch28】 Locks

### 28.6 A Failed Attempt: Just Using Loads/Stores

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 -> lock is available, 1 -> held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10         ; // spin-wait (do nothing)
11     mutex->flag = 1;          // now SET it!
12 }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

Figure 28.1: **First Attempt: A Simple Flag**

In this first attempt, the idea is quite simple: use a simple variable(flag) to indicate whether some thread has possession of a lock.

The first thread that enters the critical section will call `lock()`, which tests whether the flag is equal to 1 (in this case, it is not), and then sets the flag to 1 to indicate that the thread now holds the lock. When finished, the thread calls `unlock()` and clears the flag, thus indicating that the lock is no longer held.

If another thread happens to call `lock()` while that first thread is in the critical section, it will simply spin-wait in the while loop for that thread to call `unlock()` and clear the flag.

Once that first thread does so, the waiting thread will fall out of the while loop, set the flag to 1 for itself, and proceed into the critical section.

Unfortunately, the code has two problems: one of **correctness**, and another of **performance**.

The correctness problem is simple to see once we get used to thinking about concurrent programming. Imagine the code interleaving in the figure above; assume `flag = 0` to begin.

We can easily produce a case where **both threads set the flag to 1 and both threads are thus able to enter the critical section**. This behaviour is what professionals call “**bad**”.

The performance problem, which we will address more later on, is the fact that the way a thread waits to acquire a lock that is already held: **it endlessly checks the value of flag, a technique known as spin-waiting**.

Spin-waiting **wastes time waiting for another thread to release a lock**. The waste is exceptionally high on a uniprocessor, where the thread that the waiter is waiting for **cannot even run**.

Thread 1	Thread 2
call lock ()	
while (flag == 1)	
interrupt: switch to Thread 2	
	call lock ()
	while (flag == 1)
	flag = 1;
	interrupt: switch to Thread 1
flag = 1; // set flag to 1 (too!)	

Figure 28.2: Trace: No Mutual Exclusion