

【OS】 Day37

▼ Class	Operating System: Three Easy Pieces
📅 Date	@February 16, 2022

【Ch30】 Condition Variables

There are many cases where a thread wishes to **check whether a condition is true before continuing its execution**.

For example, a parent thread might wish to check whether a child thread has completed before continuing.

```
void *child(void *arg) {
    printf("child\n");
    // XXX how to indicate we are done?
    return NULL;
}

int main(int argc, char *argv[]) {
    printf("parent: begin\n");
    pthread_t c;
    Pthread_create(&c, NULL, child, NULL); // create child
    // XXX how to wait for child?
    printf("parent: end\n");
    return 0;
}
```

Figure 30.1: A Parent Waiting For Its Child

What we would like to see here is the following output:

```
parent: begin
child
parent: end
```

We could let the parent thread spin wait for the child thread to finish, but **this can be super inefficient**:

```

1  volatile int done = 0;
2
3  void *child(void *arg) {
4      printf("child\n");
5      done = 1;
6      return NULL;
7  }
8
9  int main(int argc, char *argv[]) {
10     printf("parent: begin\n");
11     pthread_t c;
12     Pthread_create(&c, NULL, child, NULL); // create child
13     while (done == 0)
14         ; // spin
15     printf("parent: end\n");
16     return 0;
17 }

```

Figure 30.2: Parent Waiting For Child: Spin-based Approach

30.1 Definition and Routines

To wait for a condition to become true, a thread can make use of what is known as a [condition variable](#).

A condition variable is an explicit queue that threads can [put themselves on when some state of execution](#) (i.e. some condition) [is not as desired](#) (by waiting on the condition); some other thread, when it changes said state, can then [wake one of those waiting threads](#) and thus allow them to continue.

To declare such a condition variable, one simply writes something like this:

`pthread_cond_t c;`, which declares `c` as a condition variable. A condition variable has two operations associated with it: `wait()` and `signal()`. The `wait()` call is executed when [a thread wishes to put itself to sleep](#); the `signal()` call is [executed when a thread has changed something in the program](#) and thus wants to wake a sleeping thread waiting on this condition.

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);
pthread_cond_signal(pthread_cond_t *c);
```

`wait()` call also takes a `mutex` as a parameter; it assumes that [this mutex is locked when wait\(\) is called](#).

The responsibility of `wait()` is to release the lock and put the calling thread to sleep; when the thread wakes up, it must re-acquire the lock before returning to the caller. This complexity stems from the desire to prevent certain race conditions from occurring when a thread is trying to put itself to sleep.

```
1  int done = 0;
2  pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3  pthread_cond_t c = PTHREAD_COND_INITIALIZER;
4
5  void thr_exit() {
6      Pthread_mutex_lock(&m);
7      done = 1;
8      Pthread_cond_signal(&c);
9      Pthread_mutex_unlock(&m);
10 }
11
12 void *child(void *arg) {
13     printf("child\n");
14     thr_exit();
15     return NULL;
16 }
17
18 void thr_join() {
19     Pthread_mutex_lock(&m);
20     while (done == 0)
21         Pthread_cond_wait(&c, &m);
22     Pthread_mutex_unlock(&m);
23 }
24
25 int main(int argc, char *argv[]) {
26     printf("parent: begin\n");
27     pthread_t p;
28     Pthread_create(&p, NULL, child, NULL);
29     thr_join();
30     printf("parent: end\n");
31     return 0;
32 }
```

Figure 30.3: Parent Waiting For Child: Use A Condition Variable

There are two cases to consider:

1. In the first, the parents creates the child thread but continues running itself and thus immediately calls into `thr_join()` to wait for the child thread to complete. In this case, it will acquire the lock, check if the child is done, and put itself to sleep by calling `wait()`.

The child will eventually run, print the message `"child"`, and call `thr_exit()` to wake the parent thread; this code just grabs the lock, sets the state variable `done`, and **signals the parents thus waking it**.

Finally, the parent will run, unlock the lock, and print the final message `"parent: end"`.

2. In the second case, the child runs immediately upon creation, sets `done` to 1, calls **signal to wake a sleeping thread**(but there is none, so it just returns), and is done. The parent then runs, calls `thr_join()`, sees that `done` is 1, and thus does not wait and returns.

Note: Just an if statement when deciding whether to wait on the condition will suffice our need here. It is always a good idea to use while.