

# 【OS】 Day44(3)

|         |                                     |
|---------|-------------------------------------|
| ▼ Class | Operating System: Three Easy Pieces |
| 📅 Date  | @February 28, 2022                  |

## 【Ch39】 Interlude: Files and Directories

### 39.9 Getting Information About Files

Beyond file access, we expect the file system to keep a fair amount of information about each file it is storing. We generally call such data about files [metadata](#).

To see the metadata for a certain file, we can use the `stat()` or `fstat()` system calls. These calls [take a pathname](#)(or file descriptor) to a file and fill in a `stat` structres shown below.

```
struct stat {
    dev_t      st_dev;        // ID of device containing file
    ino_t      st_ino;        // inode number
    mode_t     st_mode;       // protection
    nlink_t    st_nlink;      // number of hard links
    uid_t      st_uid;        // user ID of owner
    gid_t      st_gid;        // group ID of owner
    dev_t      st_rdev;       // device ID (if special file)
    off_t      st_size;       // total size, in bytes
    blksize_t  st_blksize;    // blocksize for filesystem I/O
    blkcnt_t   st_blocks;     // number of blocks allocated
    time_t     st_atime;      // time of last access
    time_t     st_mtime;      // time of last modification
    time_t     st_ctime;      // time of last status change
};
```

Figure 39.5: The `stat` structure.

There is a lot of information kept about each file, including its [size](#)(in bytes), its [low-level name](#)(i.e. inode name), some [ownership information](#), and some information about when the file was accessed or modified.

To see this information, we can use the command line tool `stat`.

```
prompt> echo hello > file
prompt> stat file
  File: 'file'
  Size: 6   Blocks: 8   IO Block: 4096   regular file
Device: 811h/2065d Inode: 67158084     Links: 1
Access: (0640/-rw-r-----)  Uid: (30686/remzi)
  Gid: (30686/remzi)
Access: 2011-05-03 15:50:20.157594748 -0500
Modify: 2011-05-03 15:50:20.157594748 -0500
Change: 2011-05-03 15:50:20.157594748 -0500
```

Each file system usually keeps this information in a structure called an [inode](#).

## 39.10 Removing Files

*How do we delete files?*

We can just run the program `rm`.

```
prompt> strace rm foo
...
unlink("foo");
...
```

`unlink()` takes the name of the file to be removed, and returns zero upon success.

## 39.11 Making Directories

We can **never write to a directory directly**. We can only **update a directory indirectly by creating files, directories, or other object types** within it. In this way, the file system makes sure that **directory contents are as expected**.

To create a directory, a single system call, `mkdir()`, is available. The `mkdir` program can be used **to create such a directory**.

```
prompt> strace mkdir foo
...
mkdir("foo", 0777)
```

When such a directory is created, it is **considered “empty”** although it does have a bare minimum of contents. Specifically, an empty directory has two entries: one entry that **refers to itself**, and one entry that **refers to its parent**.

The format is referred to as the **“.”(dot) directory**, and the latter as **“..”(dot-dot)** .

We can see these directories by passing a flag(**-a**) to the program **ls** :

```
prompt> ls -a
./ ../
```

## 39.12 Reading Directories

Below is an example program that **prints the contents of a directory**. The program uses three calls **opendir()** , **readdir()** , and **closedir()** .

```
int main(int argc, char *argv[]) {
    DIR *dp = opendir(".");
    assert(dp != NULL);
    struct dirent *d;
    while((d = readdir(dp)) != NULL) {
        printf("%lu %s\n", (unsigned long) d->d_ino, d->d_name);
    }
    closedir(dp);
    return 0;
}
```

The declaration below shows the information available within each directory entry in the struct dirent data structure:

```
struct dirent {
    char d_name[256]; //filename
    ino_t d_ino; //inode number
    off_t d_off; //offset to the next dirent
    unsigned short d_reclen; //length of this record
}
```

```
    unsigned char d_type; // type of file
};
```

### 39.13 Deleting Directories

Finally, we can **delete a directory** with a call to `rmdir()` (which is used by the program of the same name, `rmdir`). Unlike file deletion, **removing directories is more dangerous**, as we could potentially delete a large amount of data with a single command.

Thus, `rmdir()` has **the requirement that the directory be empty** (i.e. only has “.” and “..” entries) before it is deleted. If we try to delete a non-empty directory, the call to `rmdir()` simply will fail.