

【OS】 Day45

▼ Class	Operating System: Three Easy Pieces
📅 Date	@March 2, 2022

【Ch39】 Interlude: Files and Directories

39.14 Hard Links

We now come back to the mystery of *why removing a file is performed via `unlink()`*, by understanding a *new way to make an entry in the file system tree*, through a system call known as `link()`.

The `link()` system call *takes two arguments*, an *old pathname* and a *new one*; when we “link” a new file name to an old one, we essentially *create another way to refer to the same file*.

The command-line program `ln` is used to do this, as we see in this example:

```
prompt> echo hello > file
prompt> cat file
hello
prompt> ln file file2
prompt> cat file2
hello
```

Here we created a file with the work “hello” in it, and called the file `file`. We then create a *hard link* to that file using the `ln` program. After this, we can examine the file by either opening `file` or `file2`.

The way that `link()` works is that it simply *creates another name in the directory we are creating the link to*, and *refers it to the same inode number* of the original file.

The file is *not copied in any way*; rather, we now just have two human-readable names(`file` and `file2`) that *both refer to the same file*.

```
prompt> ls -li file file2
67158084 file
67158084 file2
```

By now we might start to see why `unlink()` is called `unlink()`. When we create a file, we are really doing two things:

1. First, we are making a structure(the inode) that will **track virtually all relevant information about the file**, including its size, where its blocks are on disk, and so forth.
2. Second, we are **linking a human-readable name to that file**, and putting that link into a directory

Thus, to **remove a file from the file system**, we call `unlink()`. In the example above, we could **remove the file named `file`** and still access the file without difficulty

```
prompt> rm file
remove 'file'
prompt> cat file2
hello
```

Every time `unlink()` is called, the OS **checks the reference count in the inode**. Only when the reference count reaches 0, the file is truly deleted.