

# 【Linux Programming】 Day4

▼ Class	Understanding Linux/Unix Programming
📅 Date	@March 15, 2022

## 【Ch2】 How does who do it?

*Read the .h files*

On most Unix machines, the header files for system information are stored in the directory called `/usr/include`. When the C compiler sees a line such as

```
#include <stdio.h>
```

it looks for that file in `/usr/include`.

Note: The `utmp.h` file on my machine is in the path `/usr/include/bits/utmp.h`

We use `more` to read the file:

```
$ more /usr/include/utmp.h
...
#define UTMP_FILE      "/var/adm/utmp"
#define WTMP_FILE      "/var/adm/wtmp"
#include <sys/types.h> /* for pid_t, time_t */

/*
 * Structure of utmp and wtmp files.
 *
 * Assuming these numbers is unwise.
 */

#define ut_name ut_user          /* compatibility */
struct utmp {
    char    ut_user[32];          /* User login name */
    char    ut_id[14];           /* /etc/inittab id- IDENT_LEN in
 * init */
    char    ut_line[32];         /* device name (console, lnxx) */
    short   ut_type;             /* type of entry */
    pid_t   ut_pid;              /* process id */
    struct exit_status {
        short   e_termination; /* Process termination status */
        short   e_exit;        /* Process exit status */
    } ut_exit;                  /* The exit status of a process
 * marked as DEAD_PROCESS.
 */
    time_t   ut_time;            /* time entry was made */
    char     ut_host[64];        /* host name same as
 * MAXHOSTNAMELEN */
};
/* Definitions for ut_type */
utmp.h (60%)
```

Log-in records, it appears, consist of eight members.

- The `ut_user` array stores the user log-in name.
- The `ut_id` array stores the device, which means the terminal from which the user connected.
- A few lines later in the struct are `ut_time` to store the log-in time and `ut_host` to store the name of the remote computer.

This struct contains other members. These do not correspond to items who displays, but they may come in handy in some situations.

### 2.4.1 We now know how who works

By reading the manual on the topics of `who` and `utmp` and by reading the header file `/usr/include/utmp.h`, we learned how `who` works.

`who` reads structures from a file. The file contains one structure for each log-in session. We know the exact layout of the structure.

The flow of information is shown in the figure below :




FIGURE 2.2  
Data flow in the `who` command.

The file is an array, so `who` must read the records and print out the information. The simplest logic would be to read and print them one by one. *Could it be that easy?*

## 2.5 Question 3: Can I write who?

There are only two tasks we need to program:

- Read structs from a file
- Display the information stored in a struct

### 2.5.1 Question: How do I read structs from a file?

We use `getc` and `fgets` to read chars and lines from a file. *What about structs of raw data?*

*Let's read the manual*

We want to find manpages about `file` and `read`. The `-k` option accepts only one keyword, let's try file first:

```
$ man -k file
```

There are a lot of topics about files. There are 537 lines printed out.

We want to search those 537 lines for lines that contain the word "read." The Unix command called `grep` prints out lines that contain a specified pattern. We use `grep` in a pipeline as follows:

```
$ man -k file | grep read
```

```
$ man -k file | grep read
_llseek (2)          - reposition read/write file offset
fileevent (n)        - Execute a script when a channel becomes readable
                      or writable
gftype (1)           - translate a generic font file for humans to read
lseek (2)            - reposition read/write file offset
macsave (1)          - Save Mac files read from standard input
read (2)             - read from a file descriptor
readprofile (1)       - a tool to read kernel profiling information
scr_dump, scr_restore, scr_init, scr_set (3) - read (write) a curses
screen from (to) a file
tee (1)              - read from standard input and write to standard
                      output and files
$
```

The most promising item in the list is `read(2)`.

Look at the manpage in section 2 about `read`:

```
$ man 2 read
READ(2)                                System calls                                READ(2)

NAME
    read - read from a file descriptor

SYNOPSIS
    #include <unistd.h>

    ssize_t read(int fd, void *buf, size_t count);

DESCRIPTION
    read() attempts to read up to count bytes from file
    descriptor fd into the buffer starting at buf.
```

## Chapter 2 Users, Files, and the Manual: who Is First

If count is zero, read() returns zero and has no other results. If count is greater than SSIZE\_MAX, the result is unspecified.

### RETURN VALUE

On success, the number of bytes read is returned (zero indicates end of file), and the file position is advanced by this number. It is not an error if this number is smaller than the number of bytes requested; this may happen for example because fewer bytes are actually available right now (maybe because we were close to end-of-file, or because we are reading from a pipe, or from a terminal), or because read() was interrupted by a signal. On error, -1 is returned, and errno is set appropriately. In this case it is left unspecified whether the file position (if any) changes.

This system calls allows us to **read a specified number of bytes from a file into a buffer**.

We want to read one struct at a time, so we can use `sizeof(struct utmp)` to **specify the number of bytes to read**.

The file reads from **a file descriptor**. *How do we get one of those?*

The page contains a reference to `open(2)`. Running

```
$ man 2 open
```

reveals how `open` works. The page for open refers to close.